

# saumit-aiml-assessment

May 12, 2024

## 1 Saumit-AIML-Assessment

### 1.1 Saumit Kunder

### 1.2 21070126078

### 1.3 AIML -B1

## 2 Libraries

```
[58]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
```

## 3 DATASET

```
[26]: train_df = pd.read_excel("./traindataset.xlsx")
```

```
[27]: train_df
```

```
[27]:
```

	ProductType	Manufacturer	Area	Code	Sourcing	Channel	Product	Size	\
0	NTM3	X1	A28		WHOLESALE			Large	
1	NTM2	X1	A9		DIRECT			Large	
2	NTM3	X2	A20		DIRECT			Large	
3	NTM3	X1	A18		WHOLESALE			Small	
4	NTM2	X1	A28		DIRECT			Large	
...	...	...	...		...		...		
550171	NTM2	X1	A5		DIRECT			Large	
550172	NTM3	X1	A14		DIRECT			Large	
550173	NTM2	X1	A5		DIRECT			Small	
550174	NTM2	X1	A7		DIRECT			Small	

550175	NTM1	X1	A3	DIRECT	Small
--------	------	----	----	--------	-------

	Product Type	Month of Sourcing	Sourcing Cost
0	Powder	2021-05-01	10.158
1	Powder	2020-10-01	134.281
2	Powder	2020-12-01	12.456
3	Powder	2021-02-01	107.220
4	Liquid	2020-11-01	197.763
...	...	...	...
550171	Powder	2020-07-01	136.469
550172	Liquid	2020-10-01	72.559
550173	Powder	2021-03-01	147.639
550174	Powder	2021-02-01	150.044
550175	Powder	2020-11-01	139.421

[550176 rows x 8 columns]

```
[28]: train_df.isnull().sum()
```

```
[28]: ProductType      0
Manufacturer      0
Area Code         0
Sourcing Channel   0
Product Size       0
Product Type       0
Month of Sourcing   0
Sourcing Cost      0
dtype: int64
```

```
[29]: train_df.dtypes
```

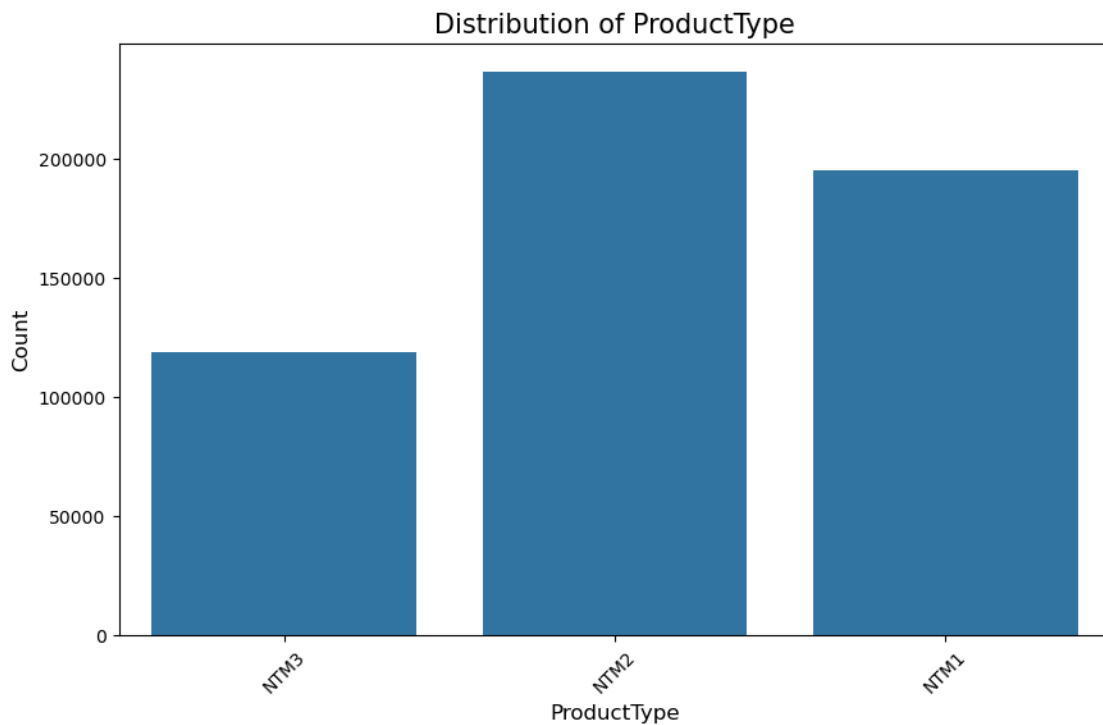
```
[29]: ProductType      object
Manufacturer      object
Area Code         object
Sourcing Channel   object
Product Size       object
Product Type       object
Month of Sourcing  datetime64[ns]
Sourcing Cost      float64
dtype: object
```


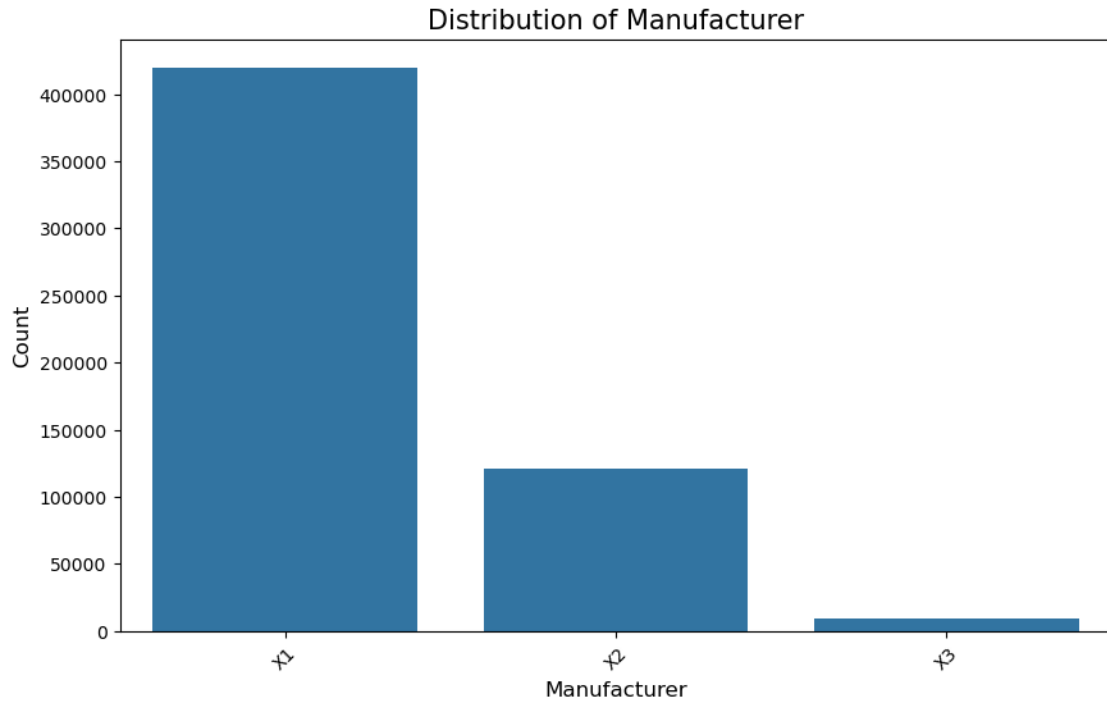
```
[30]: test_df=pd.read_excel("./testdata.xlsx")
```


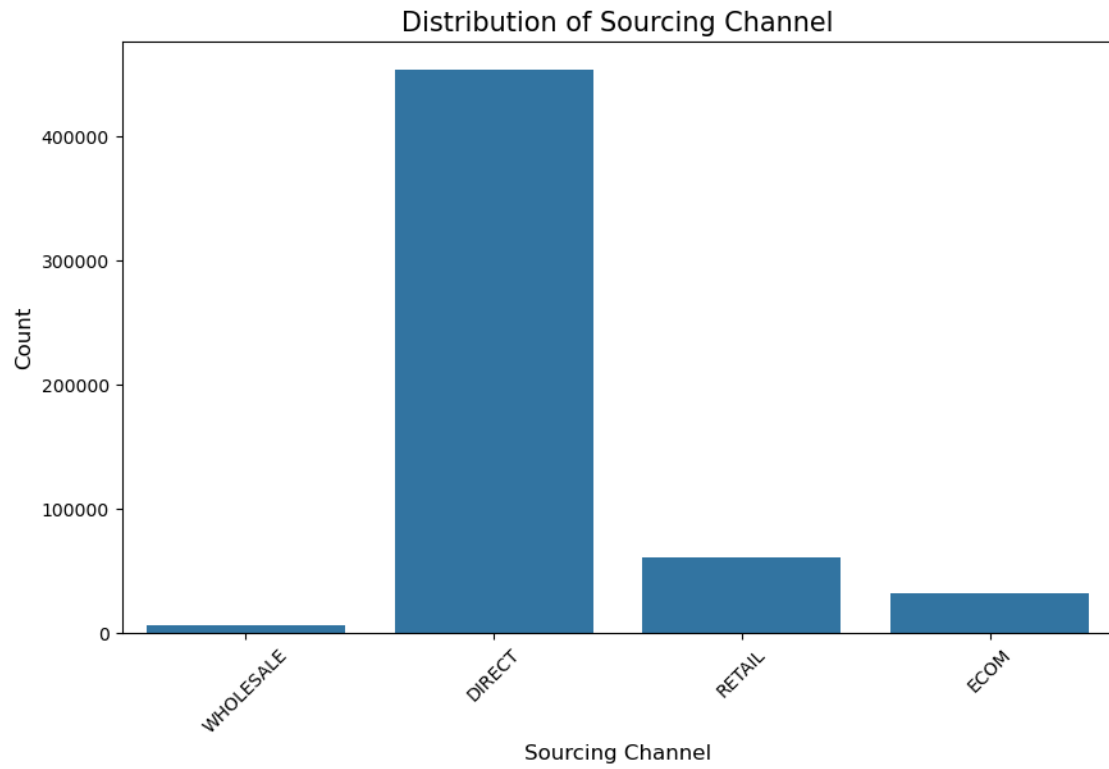
## 4 EDA

```
[39]: # List of categorical columns
categorical_columns = ['ProductType', 'Manufacturer', 'Area Code', 'Sourcing_
↳Channel', 'Product Size']

# Plot count plots for each categorical feature
for feature in categorical_columns:
    plt.figure(figsize=(10, 6))
    sns.countplot(x=feature, data=train_df)
    plt.title(f'Distribution of {feature}', fontsize=15)
    plt.xlabel(feature, fontsize=12)
    plt.ylabel('Count', fontsize=12)
    plt.xticks(rotation=45)
    plt.show()
```







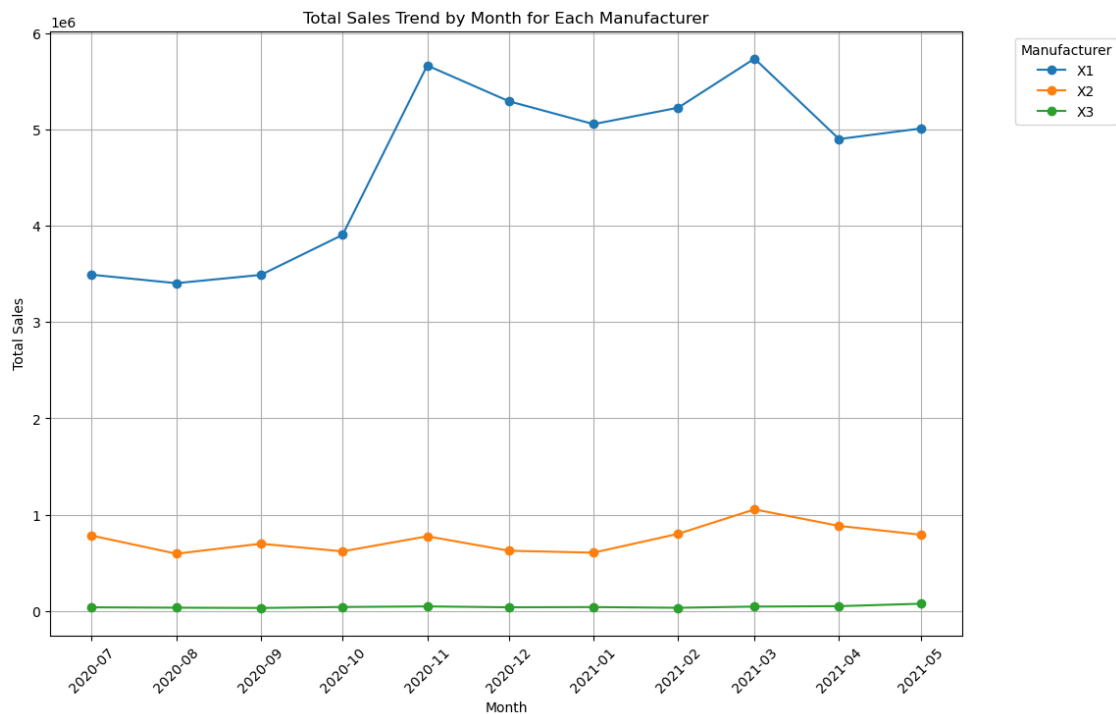
```
[31]: # Convert 'Month of Sourcing' column to datetime
train_df['Month of Sourcing'] = pd.to_datetime(train_df['Month of Sourcing'])

# Group data by both manufacturer and month and calculate total sales
monthly_sales_by_manufacturer = train_df.groupby(['Manufacturer',
↪train_df['Month of Sourcing'].dt.to_period('M')])['Sourcing Cost'].sum()

# Plotting
plt.figure(figsize=(12, 8))

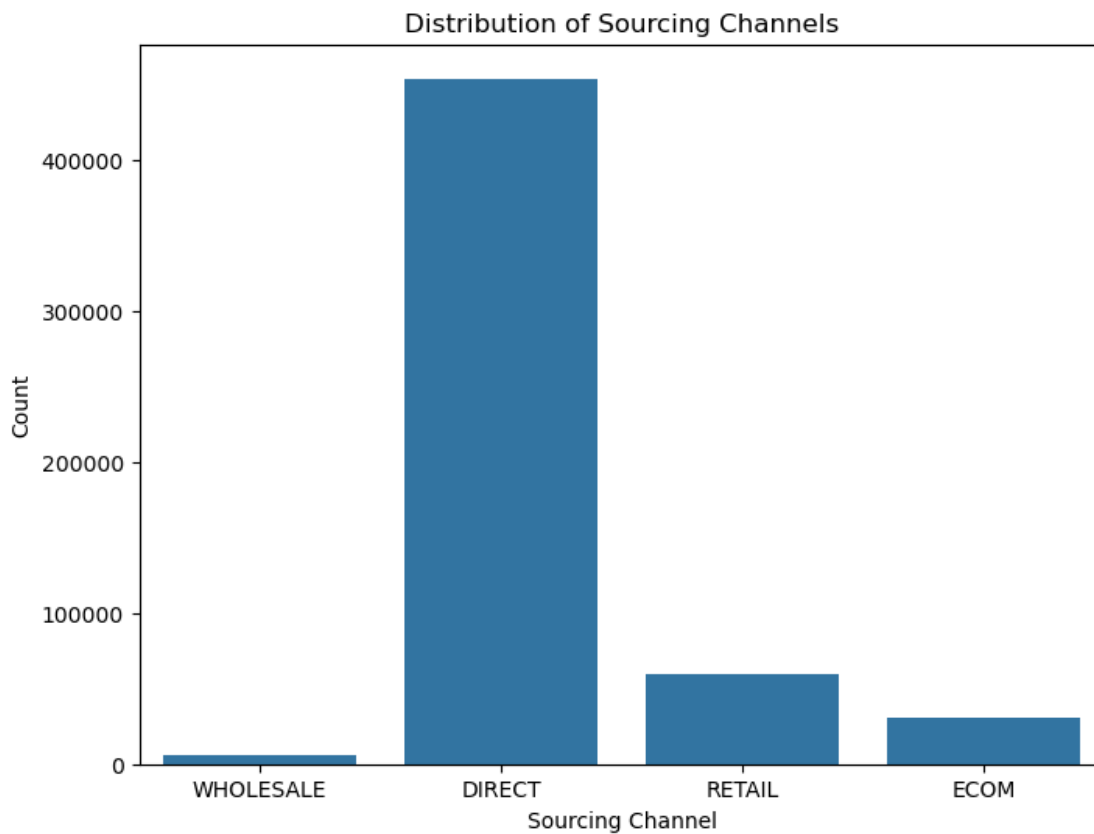
# Iterate through each manufacturer
for manufacturer, sales in monthly_sales_by_manufacturer.groupby(level=0):
    plt.plot(sales.index.get_level_values('Month of Sourcing').to_timestamp(),
↪sales.values, marker='o', linestyle='-', label=manufacturer)

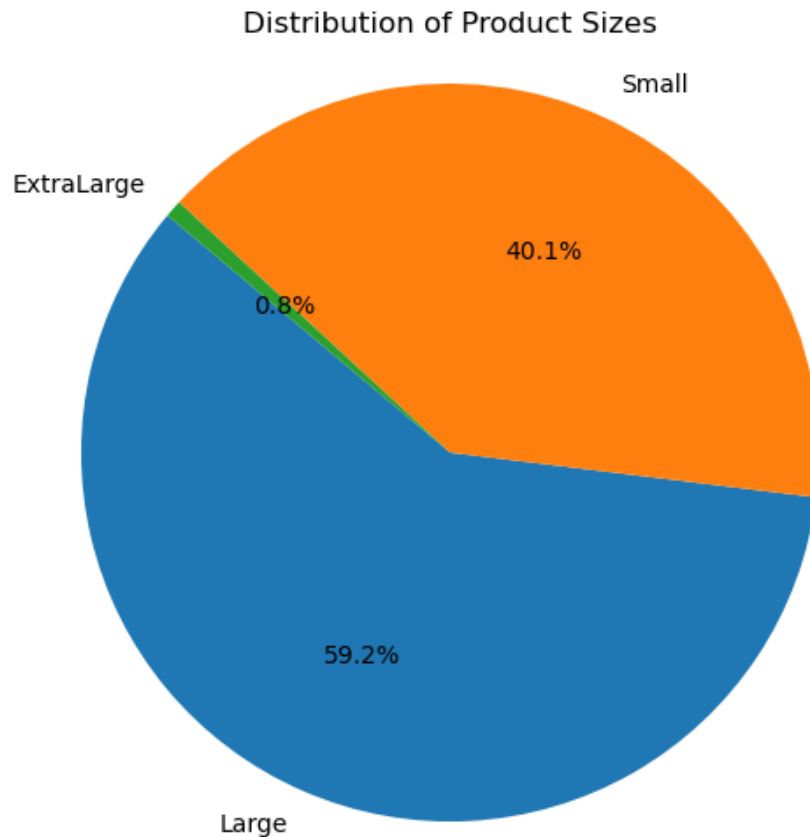
plt.title('Total Sales Trend by Month for Each Manufacturer')
plt.xlabel('Month')
plt.ylabel('Total Sales')
plt.legend(title='Manufacturer', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```



```
[32]: # Distribution of Sourcing Channels
plt.figure(figsize=(8, 6))
sns.countplot(data=train_df, x='Sourcing Channel')
plt.title('Distribution of Sourcing Channels')
plt.xlabel('Sourcing Channel')
plt.ylabel('Count')
plt.show()

# Distribution of Product Sizes
plt.figure(figsize=(8, 6))
sizes = train_df['Product Size'].value_counts()
plt.pie(sizes, labels=sizes.index, autopct='%1.1f%%', startangle=140)
plt.title('Distribution of Product Sizes')
plt.axis('equal')
plt.show()
```





```
[33]: # Calculate the frequency of each sourcing channel
channel_counts = train_df['Sourcing Channel'].value_counts()

# Plotting
plt.figure(figsize=(8, 8))
plt.pie(channel_counts, labels=channel_counts.index, autopct='%1.1f%%',
        ↪startangle=140)
plt.title('Distribution of Sourcing Channels')
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle
plt.show()

# Group data by Manufacturer, Sourcing Channel, and Product Type and calculate
↪average sourcing cost
grouped_data = train_df.groupby(['Manufacturer', 'Sourcing Channel', 'Product_
↪Type'])['Sourcing Cost'].mean().reset_index()


# Plotting
plt.figure(figsize=(12, 8))
```



```

sns.barplot(data=grouped_data, x='Manufacturer', y='Sourcing Cost',
            hue='Sourcing Channel', ci=None)
plt.title('Average Sourcing Cost by Manufacturer, Sourcing Channel, and Product Type')
plt.xlabel('Manufacturer')
plt.ylabel('Average Sourcing Cost')
plt.xticks(rotation=45)
plt.legend(title='Sourcing Channel', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

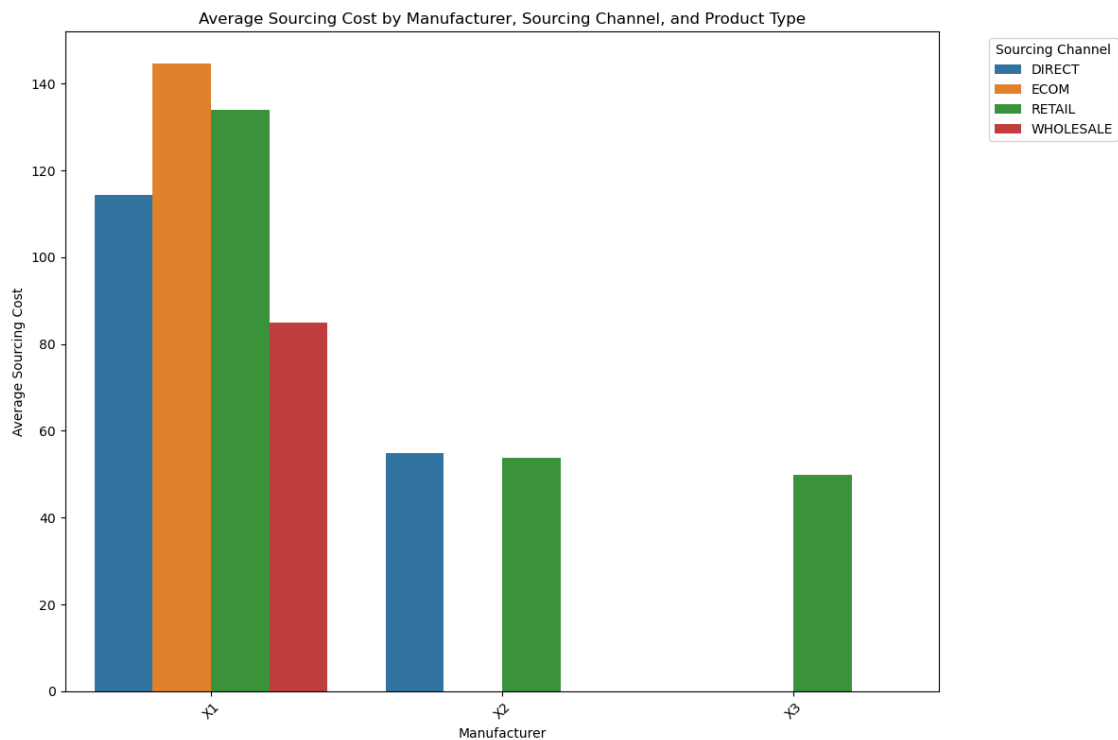
```



/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/2719704464.py:17  
: FutureWarning:

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(data=grouped_data, x='Manufacturer', y='Sourcing Cost',
hue='Sourcing Channel', ci=None)
```

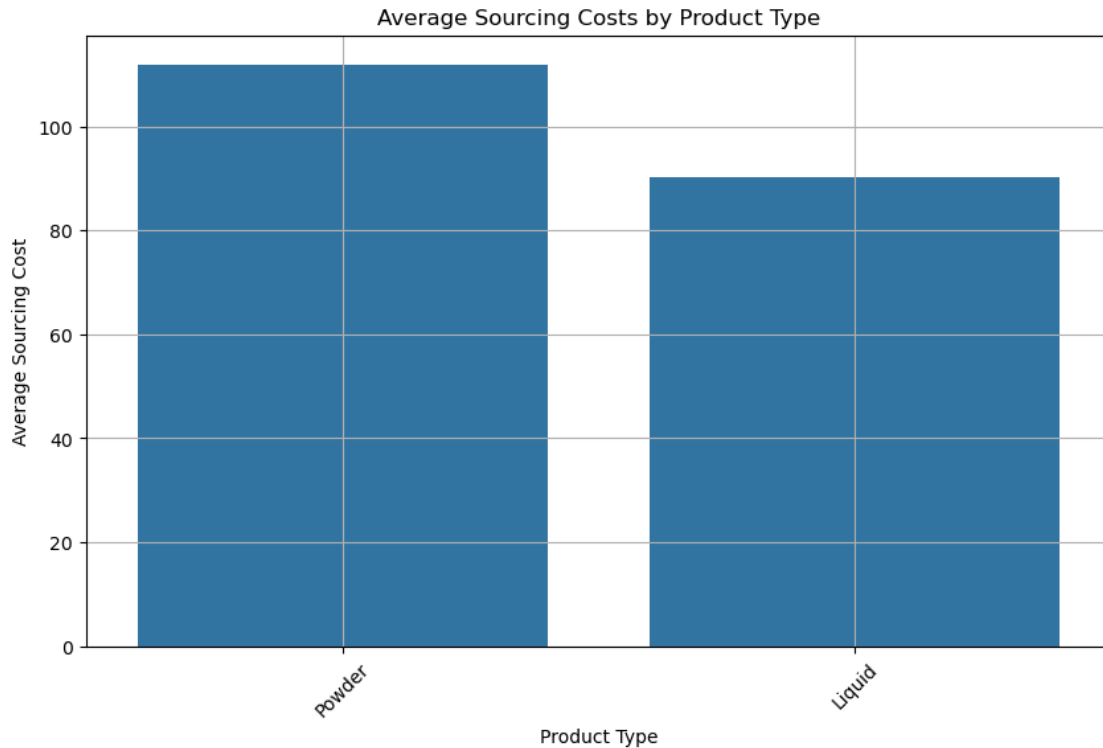


```
[34]: # Plotting
plt.figure(figsize=(10, 6))
sns.barplot(data=train_df, x='Product Type', y='Sourcing Cost', ci=None)
plt.title('Average Sourcing Costs by Product Type')
plt.xlabel('Product Type')
plt.ylabel('Average Sourcing Cost')
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/3878640288.py:3:  
FutureWarning:

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

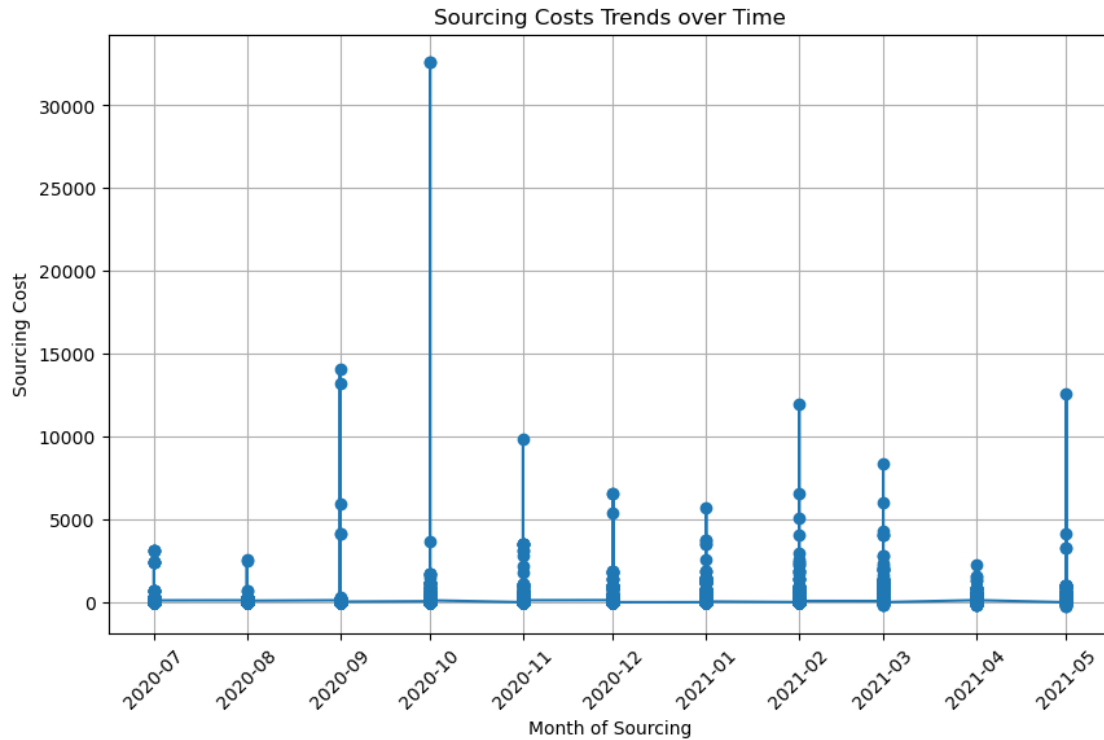
```
sns.barplot(data=train_df, x='Product Type', y='Sourcing Cost', ci=None)
```




```
[35]: # Convert 'Month of Sourcing' column to datetime
train_df['Month of Sourcing'] = pd.to_datetime(train_df['Month of Sourcing'])

# Sort the DataFrame by 'Month of Sourcing' for plotting
df = train_df.sort_values('Month of Sourcing')

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(df['Month of Sourcing'], df['Sourcing Cost'], marker='o', linestyle='-')
plt.title('Sourcing Costs Trends over Time')
plt.xlabel('Month of Sourcing')
plt.ylabel('Sourcing Cost')
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```



```
[36]: # Plotting
plt.figure(figsize=(12, 6))
sns.countplot(data=df, x='Area Code', hue='Product Type', order=df['Area Code'].
    ↳value_counts().index)
plt.title('Distribution of Transactions across Different Areas by Product Type')
plt.xlabel('Area Code')
plt.ylabel('Number of Transactions')
plt.xticks(rotation=45)
plt.legend(title='Product Type', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()
```



[ ]:

## 4.1 Outlier Detection

```
[40]: from scipy import stats

# Calculate z-scores
z_scores = stats.zscore(df['Sourcing Cost'])

# Define threshold for outlier detection
threshold = 3

# Identify outliers
outliers = df['Sourcing Cost'][abs(z_scores) > threshold]

# Print outliers
print("Outliers identified using z-score method:")
print(outliers)
# Count outliers identified using z-score method
outliers_count_zscore = len(outliers)
print("Number of outliers identified using z-score method:",
      outliers_count_zscore)
```

Outliers identified using z-score method:

```
332368    3125.792
358951    3125.792
30326      720.000
147317    2412.380
138809    2412.380
```

```

...
430963    948.896
431023    536.562
431121    948.896
72238     444.600
73142     450.047
Name: Sourcing Cost, Length: 1364, dtype: float64
Number of outliers identified using z-score method: 1364

```

```

[41]: # Calculate IQR
Q1 = df['Sourcing Cost'].quantile(0.25)
Q3 = df['Sourcing Cost'].quantile(0.75)
IQR = Q3 - Q1

# Define thresholds for outlier detection
lower_threshold = Q1 - 1.5 * IQR
upper_threshold = Q3 + 1.5 * IQR

# Identify outliers
outliers_iqr = df[(df['Sourcing Cost'] < lower_threshold) | (df['Sourcing_Cost'] > upper_threshold)]['Sourcing Cost']

# Print outliers
print("Outliers identified using IQR method:")
print(outliers_iqr)
# Count outliers identified using IQR method
outliers_count_iqr = len(outliers_iqr)
print("Number of outliers identified using IQR method:", outliers_count_iqr)

```

```

Outliers identified using IQR method:
342015    288.119
481412    288.119
41555     336.522
486424    288.119
327259    288.119
...
430963    948.896
431023    536.562
431121    948.896
72238     444.600
73142     450.047
Name: Sourcing Cost, Length: 2666, dtype: float64
Number of outliers identified using IQR method: 2666

```

## 5 Model Training

## Random Forest

```
[66]: # Convert 'Month of Sourcing' to datetime
train_df['Month of Sourcing'] = pd.to_datetime(train_df['Month of Sourcing'])
test_df['Month of Sourcing'] = pd.to_datetime(test_df['Month of Sourcing'])

# Extract features from datetime column
train_df['Year'] = train_df['Month of Sourcing'].dt.year
train_df['Month'] = train_df['Month of Sourcing'].dt.month
# Add more features as needed

test_df['Year'] = test_df['Month of Sourcing'].dt.year
test_df['Month'] = test_df['Month of Sourcing'].dt.month
# Add more features as needed

# Separate features (X) and target variable (y) for train and test datasets
X_train = train_df.drop(columns=['Sourcing Cost'])
y_train = train_df['Sourcing Cost']
X_test = test_df.drop(columns=['Sourcing Cost'])
y_test = test_df['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_train_encoded = encoder.fit_transform(X_train[categorical_columns])
X_test_encoded = encoder.transform(X_test[categorical_columns])

# Convert encoded arrays to DataFrame
X_train_encoded = pd.DataFrame(X_train_encoded, columns=encoder.
    ↳get_feature_names_out(input_features=categorical_columns))
X_test_encoded = pd.DataFrame(X_test_encoded, columns=encoder.
    ↳get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_train_final = pd.concat([X_train_encoded.reset_index(drop=True),
    ↳train_df[['Year', 'Month']].reset_index(drop=True)], axis=1)
X_test_final = pd.concat([X_test_encoded.reset_index(drop=True),
    ↳test_df[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the Random Forest model
random_forest_model = RandomForestRegressor()
random_forest_model.fit(X_train_final, y_train)

# Evaluate the model
y_pred_rf = random_forest_model.predict(X_test_final)
mse_rf = mean_squared_error(y_test, y_pred_rf)
mae_rf = mean_absolute_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)
rmse_rf = np.sqrt(mse_rf)
```

```

print("Random Forest - Mean Squared Error:", mse_rf)
print("Random Forest - Mean Absolute Error:", mae_rf)
print("Random Forest - R-squared:", r2_rf)
print("Random Forest - Root Mean Squared Error:", rmse_rf)

# Plot predictions vs. actual values
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred_rf, color='blue', alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Random Forest - Predictions vs. Actual')
plt.grid(True)
plt.show()

```

```

/Users/saumit/anaconda3/lib/python3.11/site-
packages/sklearn/preprocessing/_encoders.py:975: FutureWarning: `sparse` was
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.
`sparse_output` is ignored unless you leave `sparse` to its default value.
  warnings.warn(

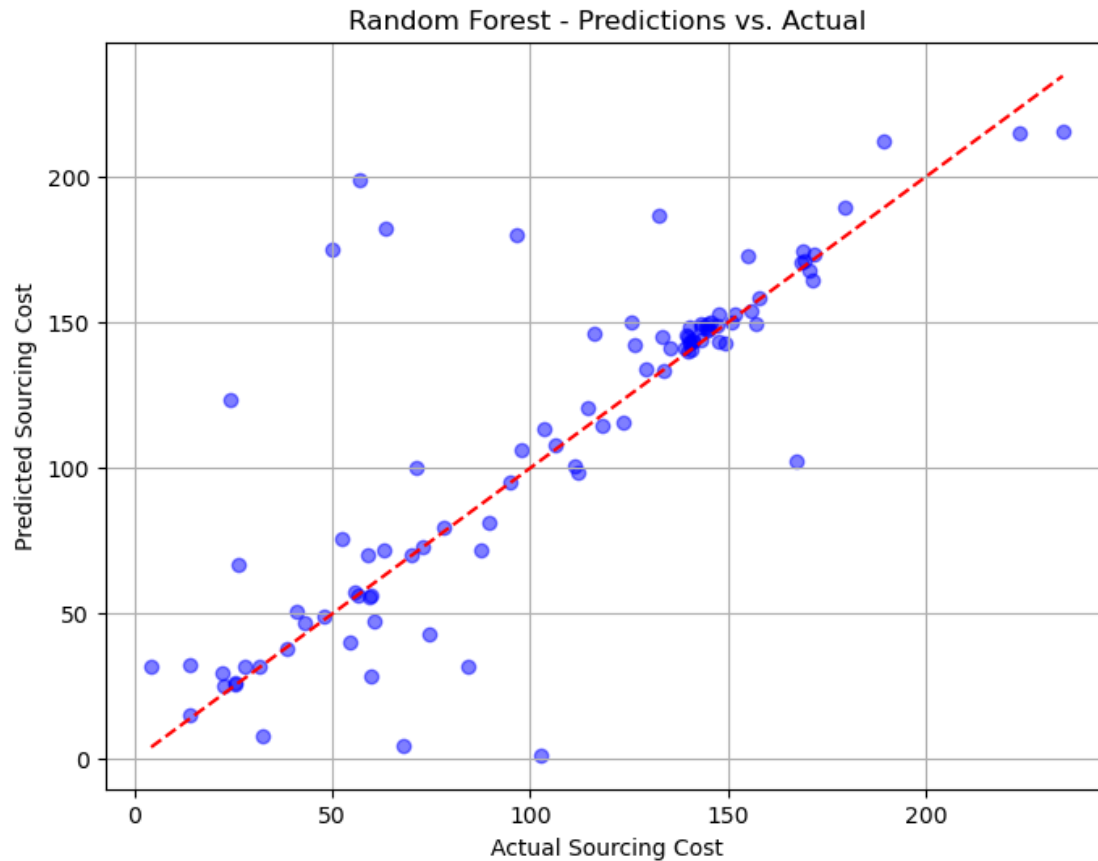
```

```

Random Forest - Mean Squared Error: 1074.392396515443
Random Forest - Mean Absolute Error: 16.569576087977335
Random Forest - R-squared: 0.6039775180717608
Random Forest - Root Mean Squared Error: 32.77792544557149

```





Random forest model using z-score for outlier/anomaly detection : Here the outliers are filtered out using z-score and then the model is trained.

```
[67]: # Filter out outliers from the original dataset
df_no_outliers = df[abs(z_scores) <= threshold]

# Identify categorical columns
categorical_columns = df_no_outliers.select_dtypes(include=['object']).columns.
    tolist()

# Convert 'Month of Sourcing' to datetime
df_no_outliers['Month of Sourcing'] = pd.to_datetime(df_no_outliers['Month of_
    Sourcing'])

# Extract features from datetime column
df_no_outliers['Year'] = df_no_outliers['Month of Sourcing'].dt.year
df_no_outliers['Month'] = df_no_outliers['Month of Sourcing'].dt.month

# Separate features (X) and target variable (y)
X = df_no_outliers.drop(columns=['Sourcing Cost'])
```

```

y = df_no_outliers['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
    ↳get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True),
    ↳df_no_outliers[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the random forest model
model = RandomForestRegressor()
model.fit(X_final, y)

# Predict using the trained model
y_pred = model.predict(X_final)

# Evaluate the model
mse_rf_zscore = mean_squared_error(y, y_pred)
mae_rf_zscore = mean_absolute_error(y, y_pred)
r2_rf_zscore = r2_score(y, y_pred)
rmse_rf_zscore = np.sqrt(mse_rf_zscore)

print("Mean Squared Error:", mse_rf_zscore)
print("Mean Absolute Error:", mae_rf_zscore)
print("R-squared:", r2_rf_zscore)
print("Root Mean Squared Error:", rmse_rf_zscore)

# Plot predictions vs. actual values
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Random Forest - Predictions vs. Actual (Without Outliers)')
plt.grid(True)
plt.show()

```


/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/413946491.py:8:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df\_no\_outliers['Month of Sourcing'] = pd.to\_datetime(df\_no\_outliers['Month of Sourcing'])  
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/413946491.py:11:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df\_no\_outliers['Year'] = df\_no\_outliers['Month of Sourcing'].dt.year  
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/413946491.py:12:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df\_no\_outliers['Month'] = df\_no\_outliers['Month of Sourcing'].dt.month  
/Users/saumit/anaconda3/lib/python3.11/site-packages/sklearn/preprocessing/\_encoders.py:975: FutureWarning: `sparse` was renamed to `sparse\_output` in version 1.2 and will be removed in 1.4.  
`sparse\_output` is ignored unless you leave `sparse` to its default value.  
warnings.warn(

Mean Squared Error: 483.1666160749332  
Mean Absolute Error: 9.808285055001443  
R-squared: 0.8515149784801571  
Root Mean Squared Error: 21.98105129594427



Random forest model using IRQ for outlier/anomaly detection : Here the outliers are filtered out using IQR(interquartile range) and then the model is trained.

```
[68]: # Filter out outliers
df_no_outliers_iqr = df[(df['Sourcing Cost'] >= lower_threshold) &
    ↪(df['Sourcing Cost'] <= upper_threshold)]

# Identify categorical columns
categorical_columns = df_no_outliers_iqr.select_dtypes(include=['object']).
    ↪columns.tolist()

# Convert 'Month of Sourcing' to datetime
df_no_outliers_iqr['Month of Sourcing'] = pd.
    ↪to_datetime(df_no_outliers_iqr['Month of Sourcing'])

# Extract features from datetime column
df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year
df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month

# Separate features (X) and target variable (y)
```

```

X = df_no_outliers_iqr.drop(columns=['Sourcing Cost'])
y = df_no_outliers_iqr['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
    ↳get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True),
    ↳df_no_outliers_iqr[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the random forest model
model = RandomForestRegressor()
model.fit(X_final, y)

# Predict using the trained model
y_pred = model.predict(X_final)

# Evaluate the model
mse_rf_iqr = mean_squared_error(y, y_pred)
mae_rf_iqr = mean_absolute_error(y, y_pred)
r2_rf_iqr = r2_score(y, y_pred)
rmse_rf_iqr = np.sqrt(mse_rf_iqr)

print("Mean Squared Error:", mse_rf_iqr)
print("Mean Absolute Error:", mae_rf_iqr)
print("R-squared:", r2_rf_iqr)
print("Root Mean Squared Error:", rmse_rf_iqr)

# Plot predictions vs. actual values
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Random Forest - Predictions vs. Actual (Without Outliers - IQR)')
plt.grid(True)
plt.show()

```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/2682994781.py:8:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers_iqr['Month of Sourcing'] =  
pd.to_datetime(df_no_outliers_iqr['Month of Sourcing'])  
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/2682994781.py:11  
: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


```
df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year  
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/2682994781.py:12  
: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month  
/Users/saumit/anaconda3/lib/python3.11/site-  
packages/sklearn/preprocessing/_encoders.py:975: FutureWarning: `sparse` was  
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.  
`sparse_output` is ignored unless you leave `sparse` to its default value.  
warnings.warn(  

```

Mean Squared Error: 364.87292187978994  
Mean Absolute Error: 9.07735567447201  
R-squared: 0.8829117663377689  
Root Mean Squared Error: 19.101647098608797



## 5.1 Linear Regression

```
[46]: # Train a Linear Regression model
linear_model = LinearRegression()
linear_model.fit(X_train_final, y_train)

# Evaluate the Linear Regression model
y_pred_linear = linear_model.predict(X_test_final)
mse_linear = mean_squared_error(y_test, y_pred_linear)
mae_linear = mean_absolute_error(y_test, y_pred_linear)
r2_linear = r2_score(y_test, y_pred_linear)
rmse_linear = np.sqrt(mse_linear)

print("Linear Regression - Mean Squared Error:", mse_linear)
print("Linear Regression - Mean Absolute Error:", mae_linear)
print("Linear Regression - R-squared:", r2_linear)
print("Linear Regression - Root Mean Squared Error:", rmse_linear)

# Plot predictions vs. actual values for Linear Regression
```


```
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred_linear, color='blue', alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Linear Regression - Predictions vs. Actual')
plt.grid(True)
plt.show()
```

Linear Regression - Mean Squared Error: 1804.6054056927085

Linear Regression - Mean Absolute Error: 26.501600052194362

Linear Regression - R-squared: 0.3348200210822406

Linear Regression - Root Mean Squared Error: 42.48064742553612



linear regression model using IRQ for outlier/anomaly detection : Here the outliers are filtered out using IQR(interquartile range) and then the model is trained.

[69]: *#Filter out outliers using IQR method*



```

df_no_outliers_iqr = df[(df['Sourcing Cost'] >= lower_threshold) &
↳(df['Sourcing Cost'] <= upper_threshold)]

# Identify categorical columns
categorical_columns = df_no_outliers_iqr.select_dtypes(include=['object']).
↳columns.tolist()

# Convert 'Month of Sourcing' to datetime
df_no_outliers_iqr['Month of Sourcing'] = pd.
↳to_datetime(df_no_outliers_iqr['Month of Sourcing'])

# Extract features from datetime column
df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year
df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month

# Separate features (X) and target variable (y)
X = df_no_outliers_iqr.drop(columns=['Sourcing Cost'])
y = df_no_outliers_iqr['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
↳get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True),
↳df_no_outliers_iqr[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the Linear Regression model
linear_model_iqr = LinearRegression()
linear_model_iqr.fit(X_final, y)

# Predict using the trained model
y_pred_linear_iqr = linear_model_iqr.predict(X_final)

# Evaluate the model
mse_linear_iqr = mean_squared_error(y, y_pred_linear_iqr)
mae_linear_iqr = mean_absolute_error(y, y_pred_linear_iqr)
r2_linear_iqr = r2_score(y, y_pred_linear_iqr)
rmse_linear_iqr = np.sqrt(mse_linear_iqr)

print("Linear Regression - Mean Squared Error (IQR):", mse_linear_iqr)
print("Linear Regression - Mean Absolute Error (IQR):", mae_linear_iqr)
print("Linear Regression - R-squared (IQR):", r2_linear_iqr)

```

```

print("Linear Regression - Root Mean Squared Error (IQR):", rmse_linear_iqr)

# Plot predictions vs. actual values for Linear Regression
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred_linear_iqr, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Linear Regression - Predictions vs. Actual (Without Outliers - IQR)')
plt.grid(True)
plt.show()

```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/1106546368.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

df_no_outliers_iqr['Month of Sourcing'] =
pd.to_datetime(df_no_outliers_iqr['Month of Sourcing'])
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/1106546368.py:11
: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/1106546368.py:12
: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```


df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month
/Users/saomit/anaconda3/lib/python3.11/site-
packages/sklearn/preprocessing/_encoders.py:975: FutureWarning: `sparse` was
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.
`sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(

```

```

Linear Regression - Mean Squared Error (IQR): 700.6212281881392
Linear Regression - Mean Absolute Error (IQR): 17.021741911005222
Linear Regression - R-squared (IQR): 0.7751696627631934
Linear Regression - Root Mean Squared Error (IQR): 26.469250616293223

```



linear regression model using z-score for outlier/anomaly detection : Here the outliers are filtered out using z-score and then the model is trained.

```
[70]: # Filter out outliers from the original dataset
df_no_outliers = df[abs(z_scores) <= threshold]

# Identify categorical columns
categorical_columns = df_no_outliers.select_dtypes(include=['object']).columns.
    ↪to list()

# Convert 'Month of Sourcing' to datetime
df_no_outliers['Month of Sourcing'] = pd.to_datetime(df_no_outliers['Month of
    ↪Sourcing'])

# Extract features from datetime column
df_no_outliers['Year'] = df_no_outliers['Month of Sourcing'].dt.year
df_no_outliers['Month'] = df_no_outliers['Month of Sourcing'].dt.month

# Separate features (X) and target variable (y)
X = df_no_outliers.drop(columns=['Sourcing Cost'])
```

```

y = df_no_outliers['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
    ↳get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True),
    ↳df_no_outliers[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the Linear Regression model
linear_model = LinearRegression()
linear_model.fit(X_final, y)

# Predict using the trained model
y_pred = linear_model.predict(X_final)

# Evaluate the model
mse_lr_zscore = mean_squared_error(y, y_pred)
mae_lr_zscore = mean_absolute_error(y, y_pred)
r2_lr_zscore = r2_score(y, y_pred)
rmse_lr_zscore = np.sqrt(mse_lr_zscore)

print("Mean Squared Error:", mse_lr_zscore)
print("Mean Absolute Error:", mae_lr_zscore)
print("R-squared:", r2_lr_zscore)
print("Root Mean Squared Error:", rmse_lr_zscore)

# Plot predictions vs. actual values
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Linear Regression - Predictions vs. Actual (Without Outliers)')
plt.grid(True)
plt.show()

```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/3890826653.py:8:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers['Month of Sourcing'] = pd.to_datetime(df_no_outliers['Month of Sourcing'])
```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/3890826653.py:11  
: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers['Year'] = df_no_outliers['Month of Sourcing'].dt.year
```

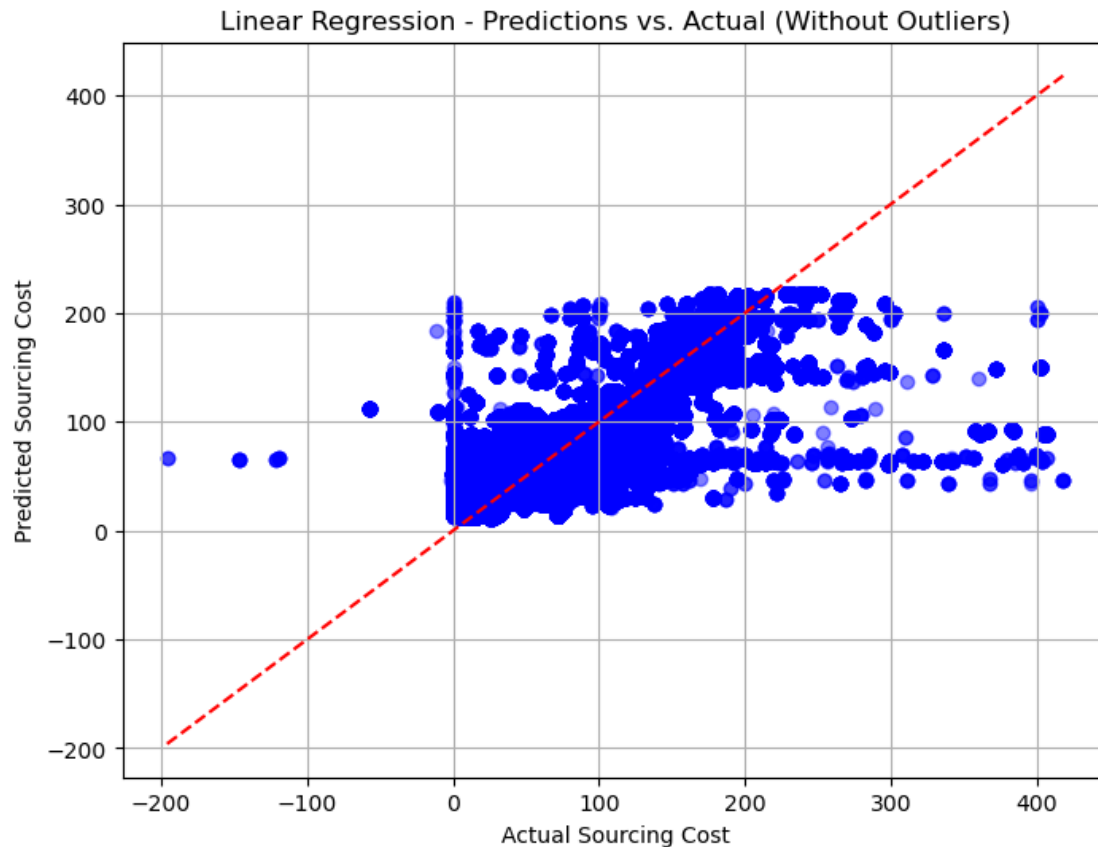
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/3890826653.py:12  
: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers['Month'] = df_no_outliers['Month of Sourcing'].dt.month
```

/Users/saumit/anaconda3/lib/python3.11/site-packages/sklearn/preprocessing/\_encoders.py:975: FutureWarning: `sparse` was renamed to `sparse\_output` in version 1.2 and will be removed in 1.4.  
`sparse\_output` is ignored unless you leave `sparse` to its default value.  
warnings.warn(

Mean Squared Error: 851.7890583295565  
Mean Absolute Error: 17.775925477507883  
R-squared: 0.7382312592623003  
Root Mean Squared Error: 29.18542544369632



## 5.2 KNN

```
[59]: # Train the KNN model
knn_model = KNeighborsRegressor()
knn_model.fit(X_train_final, y_train)

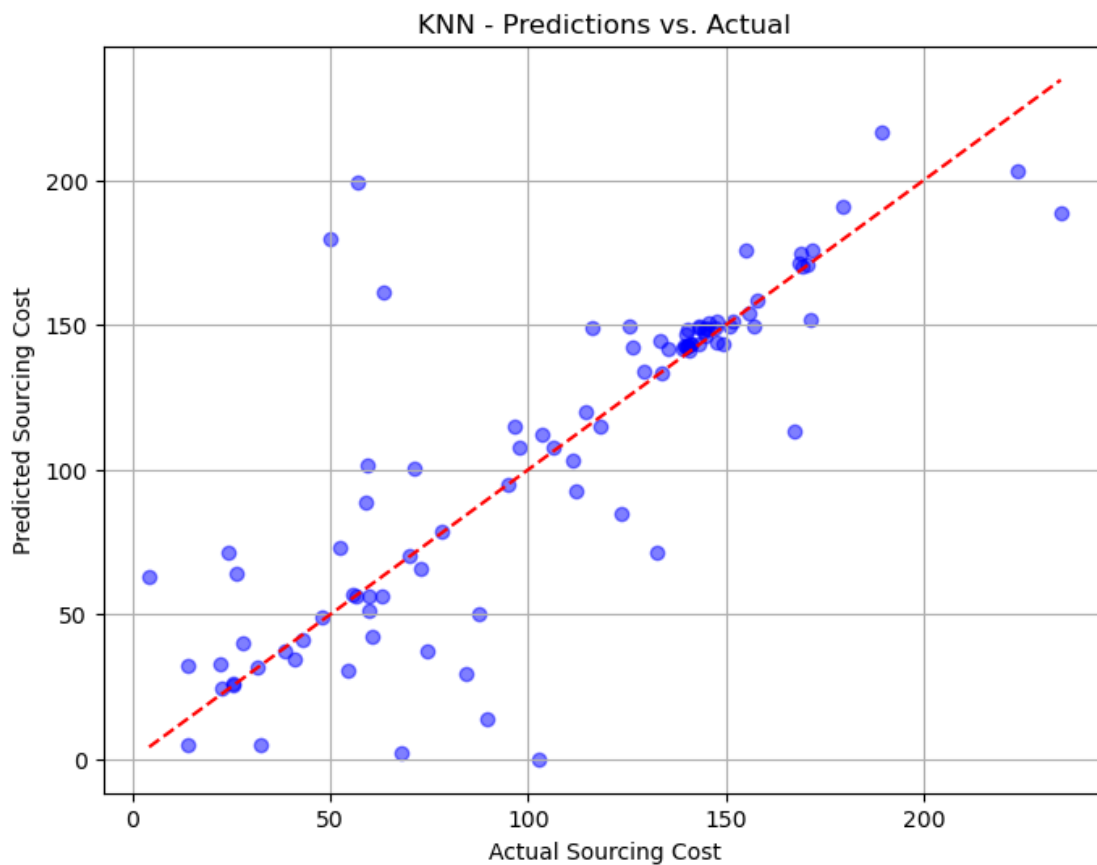
# Evaluate the KNN model
y_pred_knn = knn_model.predict(X_test_final)
mse_knn = mean_squared_error(y_test, y_pred_knn)
mae_knn = mean_absolute_error(y_test, y_pred_knn)
r2_knn = r2_score(y_test, y_pred_knn)
rmse_knn = np.sqrt(mse_knn)

print("KNN - Mean Squared Error:", mse_knn)
print("KNN - Mean Absolute Error:", mae_knn)
print("KNN - R-squared:", r2_knn)
print("KNN - Root Mean Squared Error:", rmse_knn)

# Plot predictions vs. actual values for KNN
plt.figure(figsize=(8, 6))
```

```
plt.scatter(y_test, y_pred_knn, color='blue', alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('KNN - Predictions vs. Actual')
plt.grid(True)
plt.show()
```

KNN - Mean Squared Error: 1069.2524501025  
 KNN - Mean Absolute Error: 18.1434875  
 KNN - R-squared: 0.6058721092304786  
 KNN - Root Mean Squared Error: 32.69942583750516



**KNN model using IRQ for outlier/anomaly detection :** Here the outliers are filtered out using IQR(interquartile range) and then the model is trained.

```
[60]: # Filter out outliers using IQR method
df_no_outliers_iqr = df[(df['Sourcing Cost'] >= lower_threshold) &
                        (df['Sourcing Cost'] <= upper_threshold)]
```

```

# Identify categorical columns
categorical_columns = df_no_outliers_iqr.select_dtypes(include=['object']).
    ↪columns.tolist()

# Convert 'Month of Sourcing' to datetime
df_no_outliers_iqr['Month of Sourcing'] = pd.
    ↪to_datetime(df_no_outliers_iqr['Month of Sourcing'])

# Extract features from datetime column
df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year
df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month

# Separate features (X) and target variable (y)
X = df_no_outliers_iqr.drop(columns=['Sourcing Cost'])
y = df_no_outliers_iqr['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
    ↪get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True), ↪
    ↪df_no_outliers_iqr[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the KNN model
knn_model_iqr = KNeighborsRegressor()
knn_model_iqr.fit(X_final, y)

# Predict using the trained model
y_pred_knn_iqr = knn_model_iqr.predict(X_final)

# Evaluate the model
mse_knn_iqr = mean_squared_error(y, y_pred_knn_iqr)
mae_knn_iqr = mean_absolute_error(y, y_pred_knn_iqr)
r2_knn_iqr = r2_score(y, y_pred_knn_iqr)
rmse_knn_iqr = np.sqrt(mse_knn_iqr)

print("KNN - Mean Squared Error (IQR):", mse_knn_iqr)
print("KNN - Mean Absolute Error (IQR):", mae_knn_iqr)
print("KNN - R-squared (IQR):", r2_knn_iqr)
print("KNN - Root Mean Squared Error (IQR):", rmse_knn_iqr)

```



```
# Plot predictions vs. actual values for KNN
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred_knn_iqr, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('KNN - Predictions vs. Actual (Without Outliers - IQR)')
plt.grid(True)
plt.show()
```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/3406452655.py:8:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers_iqr['Month of Sourcing'] =
pd.to_datetime(df_no_outliers_iqr['Month of Sourcing'])
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/3406452655.py:11
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/3406452655.py:12
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

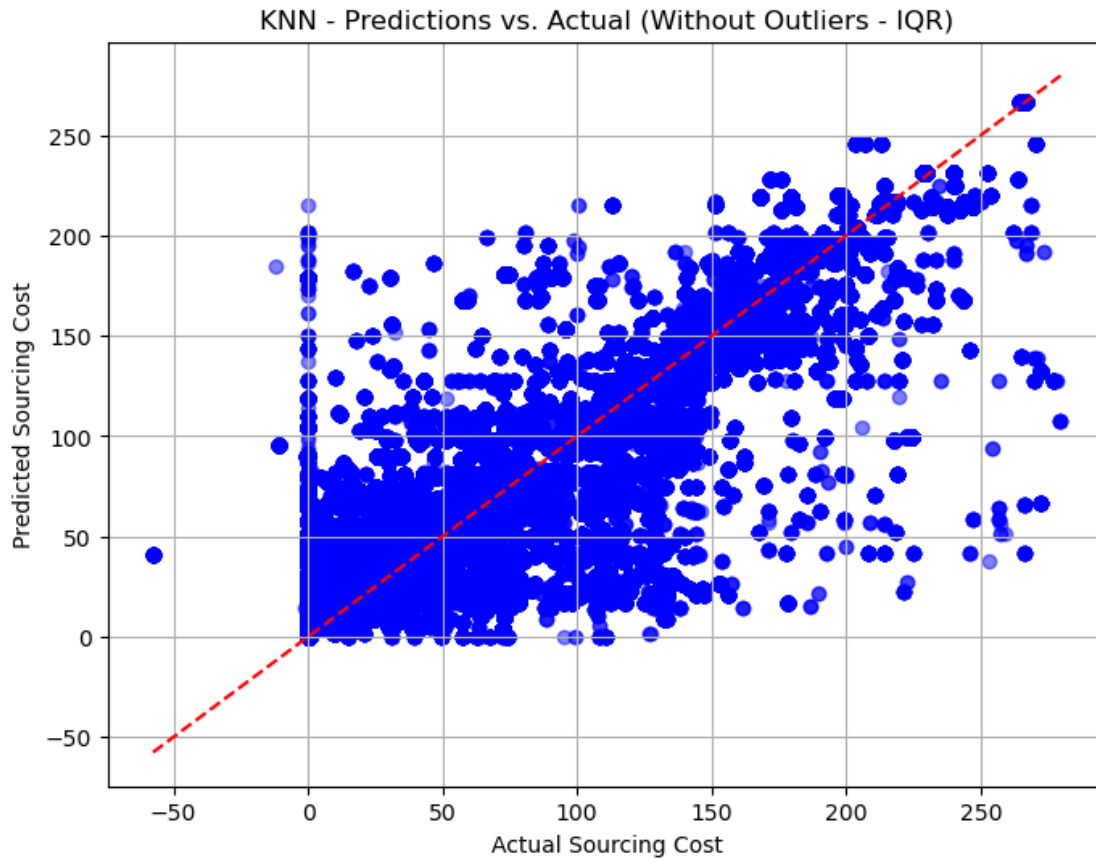
```
df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month
/Users/saumit/anaconda3/lib/python3.11/site-
packages/sklearn/preprocessing/_encoders.py:975: FutureWarning: `sparse` was
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.
`sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(
```

KNN - Mean Squared Error (IQR): 448.9442999700203

KNN - Mean Absolute Error (IQR): 9.914246489744482

KNN - R-squared (IQR): 0.8559331428997223

KNN - Root Mean Squared Error (IQR): 21.188305736184297



KNN model using z-score for outlier/anomaly detection : Here the outliers are filtered out using z-score and then the model is trained.

```
[62]: # Filter out outliers using Z-score method
df_no_outliers_zscore = df[np.abs(stats.zscore(df['Sourcing Cost'])) <
    ↳ threshold]

# Identify categorical columns
categorical_columns = df_no_outliers_zscore.select_dtypes(include=['object']).
    ↳ columns.tolist()

# Convert 'Month of Sourcing' to datetime
df_no_outliers_zscore['Month of Sourcing'] = pd.
    ↳ to_datetime(df_no_outliers_zscore['Month of Sourcing'])

# Extract features from datetime column
df_no_outliers_zscore['Year'] = df_no_outliers_zscore['Month of Sourcing'].dt.
    ↳ year
df_no_outliers_zscore['Month'] = df_no_outliers_zscore['Month of Sourcing'].dt.
    ↳ month
```

```

# Separate features (X) and target variable (y)
X = df_no_outliers_zscore.drop(columns=['Sourcing Cost'])
y = df_no_outliers_zscore['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
    ↳get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True),
    ↳df_no_outliers_zscore[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the KNN model
knn_model_zscore = KNeighborsRegressor()
knn_model_zscore.fit(X_final, y)

# Predict using the trained model
y_pred_knn_zscore = knn_model_zscore.predict(X_final)

# Evaluate the model
mse_knn_zscore = mean_squared_error(y, y_pred_knn_zscore)
mae_knn_zscore = mean_absolute_error(y, y_pred_knn_zscore)
r2_knn_zscore = r2_score(y, y_pred_knn_zscore)
rmse_knn_zscore = np.sqrt(mse_knn_zscore)

print("KNN - Mean Squared Error (Z-score):", mse_knn_zscore)
print("KNN - Mean Absolute Error (Z-score):", mae_knn_zscore)
print("KNN - R-squared (Z-score):", r2_knn_zscore)
print("KNN - Root Mean Squared Error (Z-score):", rmse_knn_zscore)

# Plot predictions vs. actual values for KNN
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred_knn_zscore, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('KNN - Predictions vs. Actual (Without Outliers - Z-score)')
plt.grid(True)
plt.show()

```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/1179194688.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers_zscore['Month of Sourcing'] =  
pd.to_datetime(df_no_outliers_zscore['Month of Sourcing'])  
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/1179194688.py:11  
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers_zscore['Year'] = df_no_outliers_zscore['Month of  
Sourcing'].dt.year  
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/1179194688.py:12  
: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

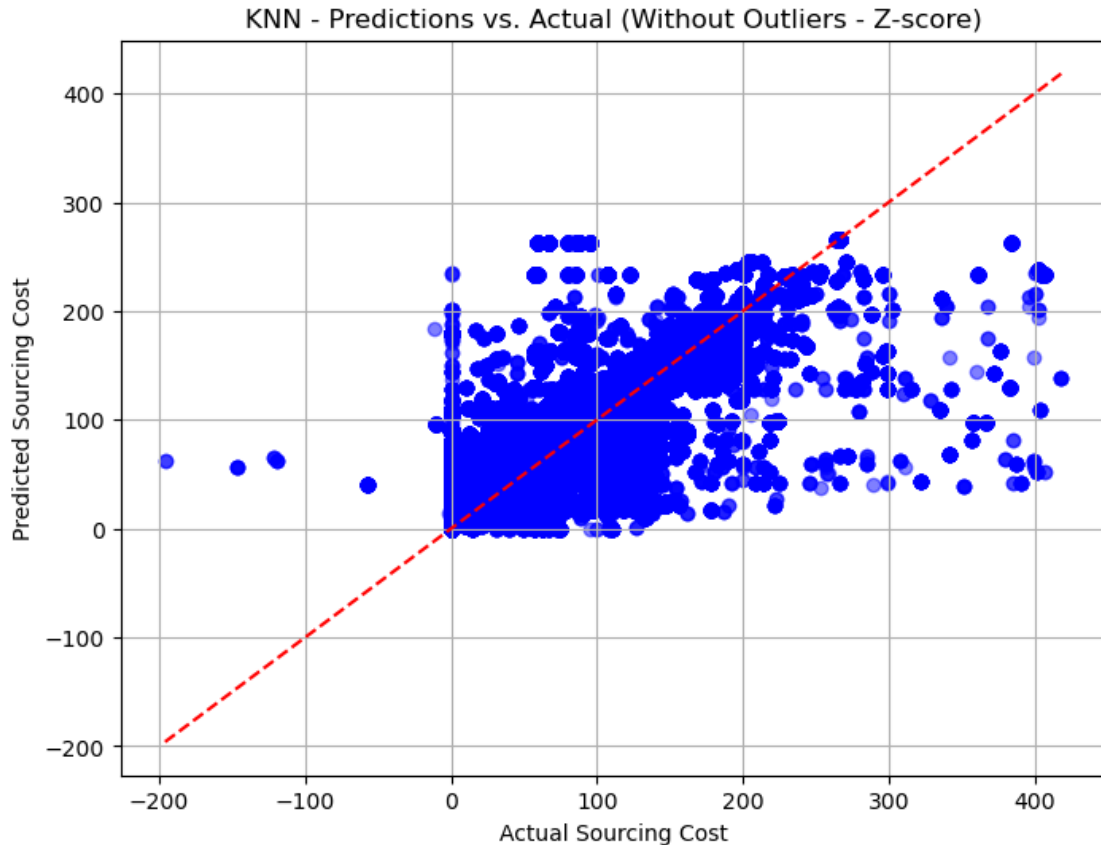
```
df_no_outliers_zscore['Month'] = df_no_outliers_zscore['Month of  
Sourcing'].dt.month  
/Users/saumit/anaconda3/lib/python3.11/site-  
packages/sklearn/preprocessing/_encoders.py:975: FutureWarning: `sparse` was  
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.  
`sparse_output` is ignored unless you leave `sparse` to its default value.  
warnings.warn(
```

KNN - Mean Squared Error (Z-score): 622.3578191839425

KNN - Mean Absolute Error (Z-score): 10.854056863916968

KNN - R-squared (Z-score): 0.8087392400466707

KNN - Root Mean Squared Error (Z-score): 24.947100416359863



[ ]:

### 5.3 Gradient Boosting

```
[63]: # Train the Gradient Boosting model
gradient_boosting_model = GradientBoostingRegressor()
gradient_boosting_model.fit(X_train_final, y_train)

# Evaluate the Gradient Boosting model
y_pred_gradient_boosting = gradient_boosting_model.predict(X_test_final)
mse_gradient_boosting = mean_squared_error(y_test, y_pred_gradient_boosting)
mae_gradient_boosting = mean_absolute_error(y_test, y_pred_gradient_boosting)
r2_gradient_boosting = r2_score(y_test, y_pred_gradient_boosting)
rmse_gradient_boosting = np.sqrt(mse_gradient_boosting)

print("Gradient Boosting - Mean Squared Error:", mse_gradient_boosting)
print("Gradient Boosting - Mean Absolute Error:", mae_gradient_boosting)
print("Gradient Boosting - R-squared:", r2_gradient_boosting)
print("Gradient Boosting - Root Mean Squared Error:", rmse_gradient_boosting)
```

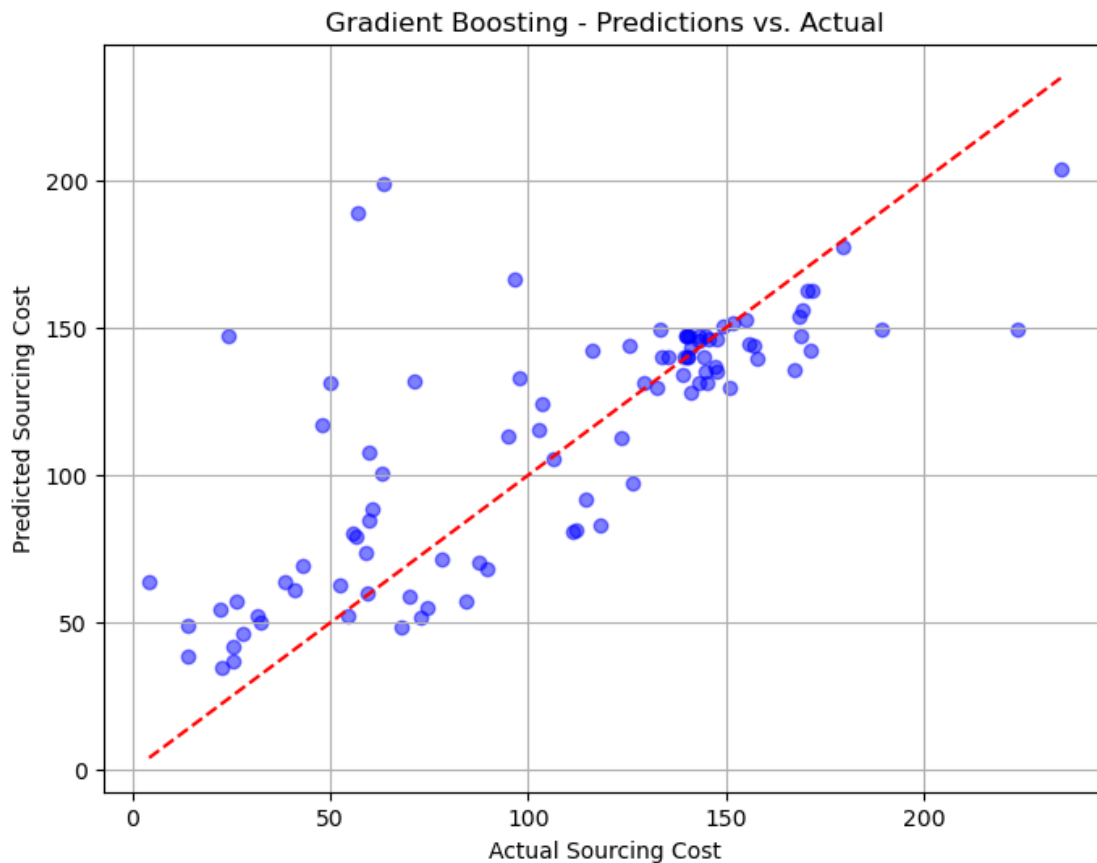
```
# Plot predictions vs. actual values for Gradient Boosting
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred_gradient_boosting, color='blue', alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Gradient Boosting - Predictions vs. Actual')
plt.grid(True)
plt.show()
```

Gradient Boosting - Mean Squared Error: 1171.4508163724956

Gradient Boosting - Mean Absolute Error: 22.597907613350674

Gradient Boosting - R-squared: 0.568201654012702

Gradient Boosting - Root Mean Squared Error: 34.22646368488126



Gradient boosting model using z-score for outlier/anomaly detection : Here the outliers are filtered out using z-score and then the model is trained.

```
[64]: # Filter out outliers using Z-score method
df_no_outliers_zscore = df[np.abs(stats.zscore(df['Sourcing Cost'])) <
    ↳threshold]

# Identify categorical columns
categorical_columns = df_no_outliers_zscore.select_dtypes(include=['object']).
    ↳columns.tolist()

# Convert 'Month of Sourcing' to datetime
df_no_outliers_zscore['Month of Sourcing'] = pd.
    ↳to_datetime(df_no_outliers_zscore['Month of Sourcing'])

# Extract features from datetime column
df_no_outliers_zscore['Year'] = df_no_outliers_zscore['Month of Sourcing'].dt.
    ↳year
df_no_outliers_zscore['Month'] = df_no_outliers_zscore['Month of Sourcing'].dt.
    ↳month

# Separate features (X) and target variable (y)
X = df_no_outliers_zscore.drop(columns=['Sourcing Cost'])
y = df_no_outliers_zscore['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
    ↳get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True),
    ↳df_no_outliers_zscore[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the Gradient Boosting model
gradient_boosting_model_zscore = GradientBoostingRegressor()
gradient_boosting_model_zscore.fit(X_final, y)

# Predict using the trained model
y_pred_gradient_boosting_zscore = gradient_boosting_model_zscore.
    ↳predict(X_final)

# Evaluate the model
mse_gradient_boosting_zscore = mean_squared_error(y,
    ↳y_pred_gradient_boosting_zscore)
```

```

mae_gradient_boosting_zscore = mean_absolute_error(y,
    ↪ y_pred_gradient_boosting_zscore)
r2_gradient_boosting_zscore = r2_score(y, y_pred_gradient_boosting_zscore)
rmse_gradient_boosting_zscore = np.sqrt(mse_gradient_boosting_zscore)

print("Gradient Boosting - Mean Squared Error (Z-score):",
    ↪ mse_gradient_boosting_zscore)
print("Gradient Boosting - Mean Absolute Error (Z-score):",
    ↪ mae_gradient_boosting_zscore)
print("Gradient Boosting - R-squared (Z-score):", r2_gradient_boosting_zscore)
print("Gradient Boosting - Root Mean Squared Error (Z-score):",
    ↪ rmse_gradient_boosting_zscore)

# Plot predictions vs. actual values for Gradient Boosting
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred_gradient_boosting_zscore, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Gradient Boosting - Predictions vs. Actual (Without Outliers -
    ↪ Z-score)')
plt.grid(True)
plt.show()

```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/2667107576.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

df_no_outliers_zscore['Month of Sourcing'] =
pd.to_datetime(df_no_outliers_zscore['Month of Sourcing'])
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/2667107576.py:11
: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

df_no_outliers_zscore['Year'] = df_no_outliers_zscore['Month of
Sourcing'].dt.year
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/2667107576.py:12
: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead




See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_no_outliers_zscore['Month'] = df_no_outliers_zscore['Month of Sourcing'].dt.month
```

/Users/saumit/anaconda3/lib/python3.11/site-packages/sklearn/preprocessing/\_encoders.py:975: FutureWarning: `sparse` was renamed to `sparse\_output` in version 1.2 and will be removed in 1.4. `sparse\_output` is ignored unless you leave `sparse` to its default value.

```
warnings.warn(
```

Gradient Boosting - Mean Squared Error (Z-score): 842.8886407385451  
 Gradient Boosting - Mean Absolute Error (Z-score): 18.707636144120055  
 Gradient Boosting - R-squared (Z-score): 0.7409665034897945  
 Gradient Boosting - Root Mean Squared Error (Z-score): 29.032544510231016



Gradient Boosting model using IRQ for outlier/anomaly detection : Here the outliers are filtered out using IRQ(interquartile range) and then the model is trained.

```
[65]: # Filter out outliers using IQR method
df_no_outliers_iqr = df[(df['Sourcing Cost'] >= lower_threshold) &
    ↪ (df['Sourcing Cost'] <= upper_threshold)]
```

```

# Identify categorical columns
categorical_columns = df_no_outliers_iqr.select_dtypes(include=['object']).
    ↪columns.tolist()

# Convert 'Month of Sourcing' to datetime
df_no_outliers_iqr['Month of Sourcing'] = pd.
    ↪to_datetime(df_no_outliers_iqr['Month of Sourcing'])

# Extract features from datetime column
df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year
df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month

# Separate features (X) and target variable (y)
X = df_no_outliers_iqr.drop(columns=['Sourcing Cost'])
y = df_no_outliers_iqr['Sourcing Cost']

# Encode categorical variables
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = encoder.fit_transform(X[categorical_columns])

# Convert encoded arrays to DataFrame
X_encoded_df = pd.DataFrame(X_encoded, columns=encoder.
    ↪get_feature_names_out(input_features=categorical_columns))

# Combine encoded categorical columns and extracted features
X_final = pd.concat([X_encoded_df.reset_index(drop=True), ↪
    ↪df_no_outliers_iqr[['Year', 'Month']].reset_index(drop=True)], axis=1)

# Train the Gradient Boosting model
gradient_boosting_model_iqr = GradientBoostingRegressor()
gradient_boosting_model_iqr.fit(X_final, y)

# Predict using the trained model
y_pred_gradient_boosting_iqr = gradient_boosting_model_iqr.predict(X_final)

# Evaluate the model
mse_gradient_boosting_iqr = mean_squared_error(y, y_pred_gradient_boosting_iqr)
mae_gradient_boosting_iqr = mean_absolute_error(y, y_pred_gradient_boosting_iqr)
r2_gradient_boosting_iqr = r2_score(y, y_pred_gradient_boosting_iqr)
rmse_gradient_boosting_iqr = np.sqrt(mse_gradient_boosting_iqr)

print("Gradient Boosting - Mean Squared Error (IQR):", ↪
    ↪mse_gradient_boosting_iqr)
print("Gradient Boosting - Mean Absolute Error (IQR):", ↪
    ↪mae_gradient_boosting_iqr)

```

```

print("Gradient Boosting - R-squared (IQR):", r2_gradient_boosting_iqr)
print("Gradient Boosting - Root Mean Squared Error (IQR):",
      rmse_gradient_boosting_iqr)

# Plot predictions vs. actual values for Gradient Boosting
plt.figure(figsize=(8, 6))
plt.scatter(y, y_pred_gradient_boosting_iqr, color='blue', alpha=0.5)
plt.plot([min(y), max(y)], [min(y), max(y)], color='red', linestyle='--')
plt.xlabel('Actual Sourcing Cost')
plt.ylabel('Predicted Sourcing Cost')
plt.title('Gradient Boosting - Predictions vs. Actual (Without Outliers - IQR)')
plt.grid(True)
plt.show()

```

/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel\_6328/1483575717.py:8:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

df_no_outliers_iqr['Month of Sourcing'] =
pd.to_datetime(df_no_outliers_iqr['Month of Sourcing'])
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/1483575717.py:11
: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

df_no_outliers_iqr['Year'] = df_no_outliers_iqr['Month of Sourcing'].dt.year
/var/folders/jw/xdjgvbw534x94g51k1ltxy640000gn/T/ipykernel_6328/1483575717.py:12
: SettingWithCopyWarning:

```

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

df_no_outliers_iqr['Month'] = df_no_outliers_iqr['Month of Sourcing'].dt.month
/Users/saumit/anaconda3/lib/python3.11/site-
packages/sklearn/preprocessing/_encoders.py:975: FutureWarning: `sparse` was
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.
`sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(


```

Gradient Boosting - Mean Squared Error (IQR): 723.8604950081095

Gradient Boosting - Mean Absolute Error (IQR): 18.285919315944867

Gradient Boosting - R-squared (IQR): 0.7677121493650024

Gradient Boosting - Root Mean Squared Error (IQR): 26.904655638162506



## 5.4 Evaluation Table

```
[71]: #Define evaluation scores for each model and outlier removal method
evaluation_data = {
    'Model': ['Random Forest (IQR)', 'Random Forest (Z-score)', 'KNN (IQR)',
    ↪ 'KNN (Z-score)',
            'Linear Regression (IQR)', 'Linear Regression (Z-score)',
            'Gradient Boosting (IQR)', 'Gradient Boosting (Z-score)'],
    'Mean Squared Error': [mse_rf_iqr, mse_rf_zscore, mse_knn_iqr,
    ↪ mse_knn_zscore,
                        mse_linear_iqr, mse_lr_zscore,
    ↪ mse_gradient_boosting_iqr, mse_gradient_boosting_zscore],
    'Mean Absolute Error': [mae_rf_iqr, mae_rf_zscore, mae_knn_iqr,
    ↪ mae_knn_zscore,
                        mae_linear_iqr, mae_lr_zscore,
    ↪ mae_gradient_boosting_iqr, mae_gradient_boosting_zscore],
    'R-squared': [r2_rf_iqr, r2_rf_zscore, r2_knn_iqr, r2_knn_zscore,
```

```

        r2_linear_iqr, r2_lr_zscore, r2_gradient_boosting_iqr, \
        r2_gradient_boosting_zscore],
        'Root Mean Squared Error': [rmse_rf_iqr, rmse_rf_zscore, rmse_knn_iqr, \
        rmse_knn_zscore,
                                     rmse_linear_iqr, rmse_lr_zscore, \
        rmse_gradient_boosting_iqr, rmse_gradient_boosting_zscore]
    }

    # Create DataFrame to store evaluation scores
    evaluation_table = pd.DataFrame(evaluation_data)

    # Display the table
    print("Evaluation Scores for Different Models with Outliers Removed (IQR and \
    Z-score):")
    print(evaluation_table)

```

Evaluation Scores for Different Models with Outliers Removed (IQR and Z-score):

	Model	Mean Squared Error	Mean Absolute Error \
0	Random Forest (IQR)	364.872922	9.077356
1	Random Forest (Z-score)	483.166616	9.808285
2	KNN (IQR)	448.944300	9.914246
3	KNN (Z-score)	622.357819	10.854057
4	Linear Regression (IQR)	700.621228	17.021742
5	Linear Regression (Z-score)	851.789058	17.775925
6	Gradient Boosting (IQR)	723.860495	18.285919
7	Gradient Boosting (Z-score)	842.888641	18.707636

	R-squared	Root Mean Squared Error
0	0.882912	19.101647
1	0.851515	21.981051
2	0.855933	21.188306
3	0.808739	24.947100
4	0.775170	26.469251
5	0.738231	29.185425
6	0.767712	26.904656
7	0.740967	29.032545

```

[72]: from tabulate import tabulate

    # Display the table using tabulate
    print(tabulate(evaluation_table, headers='keys', tablefmt='grid'))

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|      | Model |      | Mean Squared Error | Mean Absolute Error |
+====+=====+=====+=====+=====+=====+=====+=====+=====+=====+

```

==+=====+				
0	Random Forest (IQR)		364.873	
9.07736	0.882912		19.1016	
+---+-----+				
--+-----+				
1	Random Forest (Z-score)		483.167	
9.80829	0.851515		21.9811	
+---+-----+				
--+-----+				
2	KNN (IQR)		448.944	
9.91425	0.855933		21.1883	
+---+-----+				
--+-----+				
3	KNN (Z-score)		622.358	
0.808739		24.9471		10.8541
+---+-----+				
--+-----+				
4	Linear Regression (IQR)		700.621	
0.77517		26.4693		17.0217
+---+-----+				
--+-----+				
5	Linear Regression (Z-score)		851.789	
0.738231		29.1854		17.7759
+---+-----+				
--+-----+				
6	Gradient Boosting (IQR)		723.86	
0.767712		26.9047		18.2859
+---+-----+				
--+-----+				
7	Gradient Boosting (Z-score)		842.889	
0.740967		29.0325		18.7076
+---+-----+				
--+-----+				

In conclusion, Random Forest with Interquartile Range (IQR) outlier removal emerges as the top-performing model among those tested.

**Robustness to Outliers:** The IQR method effectively filters out outliers, ensuring that the Random Forest model is trained on a more representative and less biased dataset. This robustness to outliers enhances the model's ability to generalize well to unseen data.

**Ensemble Learning:** Random Forest leverages ensemble learning, combining multiple decision trees trained on random subsets of the data. This ensemble approach mitigates overfitting and improves predictive accuracy by capturing diverse patterns in the data.

**Non-linear Relationships:** Random Forest can model complex, non-linear relationships between features and the target variable. This flexibility allows the model to capture intricate patterns that may exist in the data, leading to superior predictive performance.

**Feature Importance:** Random Forest provides insights into feature importance, enabling us to

identify the most influential predictors. This feature analysis aids in understanding the underlying factors driving the predicted outcomes and can inform decision-making in real-world applications.

[ ]: